

ECN2160: Introduction

Alexandre Brilhante
Karim Itani

28 janvier 2018

0.1 Installer R

<https://www.r-project.org/>

0.2 Installer Jupyter

<http://jupyter.org/>

0.3 Installer IRKernel

<https://irkernel.github.io/>

Sur un terminal: jupyter notebook Un notebook comme celui-ci s'ouvrira sur votre fureteur.

0.4 Alternative: installer RStudio

<https://www.rstudio.com/>

0.5 Opérations de base

```
In [1]: x <- 345
```

```
x
```

```
345
```

```
In [2]: log_x <- log2(x)
```

```
log_x
```

```
8.43045255166553
```

```
In [3]: y <- 45
```

```
x-y
```

```
300
```

```
In [4]: z <- 300
```

```
  if (x-y == z) {  
    return(TRUE)  
  } else {  
    return(FALSE)  
  }
```

TRUE

```
In [5]: z <- x-x
```

z

0

```
In [6]: a <- 24
```

```
  b <- 3
```

```
  mult <- function(a, b) {  
    result <- 0  
    for (i in 1:b) {  
      result <- result+a  
    }  
    return(result)  
  }
```

```
  mult(a, b)
```

72

```
In [7]: a*b
```

72

0.6 Série chronologique

```
In [8]: # On génère 12 valeurs aléatoires.
```

```
  input <- rnorm(12)
```

```
In [9]: input
```

```
  1. 0.63268462382664 2. -0.186421772308832 3. 1.01887968323517 4. -0.297304114335721  
  5. -0.483965435795269 6. -0.33327251542413 7. 1.51388337749175 8. 0.225710957771721  
  9. -1.25736090221971 10. -1.40500551148294 11. -0.97308907337095 12. 0.464128996757302
```


```
In [10]: # On convertit nos valeurs en série chronologique.
```

```
  input <- ts(input, start = c(2018, 1), frequency = 12)
```

```
In [11]: input
```

	Jan	Feb	Mar	Apr	May	Jun
2018	0.6326846	-0.1864218	1.0188797	-0.2973041	-0.4839654	-0.3332725
	Jul	Aug	Sep	Oct	Nov	Dec
2018	1.5138834	0.2257110	-1.2573609	-1.4050055	-0.9730891	0.4641290

```
In [12]: plot.ts(input)
```




0.7 Nuage de point

```
In [25]: # On génère 10 points aléatoires.
```

```
x <- rnorm(10)
y <- rnorm(10)
```

```
In [26]: plot(x, y)
```



0.8 Régression linéaire

```
In [27]: # Création de données.
set.seed(123)
```

```
x <- 1:20 + rnorm(100, sd = 3)
z <- 1:20/4 + rnorm(100, sd = 2)
y <- -2*x + x*z/5 + 3 + rnorm(100, sd = 4)
```

```
# On crée un dataframe avec nos données.
dat <- data.frame(x = x, y = y, z = z)
```

```
head(dat)
tail(dat)
```

	x	y	z
	-0.6814269	13.317660	-1.1708131
	1.3094675	5.896216	1.0137674
	7.6761249	-13.018866	0.2566162
	4.2115252	-2.993443	0.3049148
	5.3878632	-10.136997	-0.6532371
	11.1451950	-18.052556	1.4099446
	x	y	z
95	19.08196	-22.847579	1.128397
96	14.19922	-2.428712	7.994427
97	23.56200	-10.967331	5.451418
98	22.59783	-38.571634	1.997457
99	18.29290	-20.595608	3.527668
100	16.92074	-16.944758	2.629040

Corrélation

```
In [16]: cor(dat$x, dat$y)

-0.789499506110698
```

Régression linéaire simple

```
In [17]: fit <- lm(dat$y ~ dat$x)

fit
```

Call:

```
lm(formula = dat$y ~ dat$x)
```

Coefficients:

```
(Intercept)      dat$x
      2.365      -1.256
```

Cela veut donc dire que $y = 2.365 - 1.256x$.

```
In [18]: summary(fit)
```

```

Call:
lm(formula = dat$y ~ dat$x)


Residuals:
    Min       1Q   Median       3Q      Max
-15.9923  -4.9756  -0.0181   3.9530  19.4712


Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  2.36489    1.24670   1.897  0.0608 .
dat$x       -1.25632    0.09866 -12.734 <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Residual standard error: 6.52 on 98 degrees of freedom
Multiple R-squared:  0.6233, Adjusted R-squared:  0.6195
F-statistic: 162.2 on 1 and 98 DF,  p-value: < 2.2e-16


```

```
In [19]: plot(fit)
```









Régression linéaire multiple

```
In [20]: fit <- lm(dat$z ~ dat$x + dat$y)
```

```
fit
```

Call:

```
lm(formula = dat$z ~ dat$x + dat$y)
```

Coefficients:

(Intercept)	dat\$x	dat\$y
-0.1735	0.4523	0.2049

Cela veut donc dire que $z = -0.1735 + 0.4523x + 0.2049y$.

```
In [21]: summary(fit)
```

Call:

```
lm(formula = dat$z ~ dat$x + dat$y)
```

Residuals:

Min	1Q	Median	3Q	Max
-4.1060	-0.8619	-0.1232	1.0091	5.6455

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.17355	0.30721	-0.565	0.573
dat\$x	0.45229	0.03890	11.626	< 2e-16 ***
dat\$y	0.20491	0.02445	8.382	4.12e-13 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.578 on 97 degrees of freedom

Multiple R-squared: 0.5852, Adjusted R-squared: 0.5767

F-statistic: 68.43 on 2 and 97 DF, p-value: < 2.2e-16

```
In [22]: plot(fit)
```

